

# DANIEL DUCLOS-CAVALCANTI

## Computer Engineer

New York, New York

XXX-XXX-XXXX | U.S. Citizen | [daniel@duclos.dev](mailto:daniel@duclos.dev) | [www.duclos.dev](http://www.duclos.dev) | [linkedin/duclos-cavalcanti](https://linkedin/duclos-cavalcanti) | [github/duclos-cavalcanti](https://github/duclos-cavalcanti)

## Education

### Technical University of Munich

Oct 2020 – Oct 2024

*M.Sc. Electrical and Computer Engineering*

*Munich, Germany*

- Visiting Non-Degree Graduate Student: **New York University** – **GPA 4.0** (2023 – 2024)
- Master Thesis: **"VM Selection Heuristic for Financial Exchanges in the Cloud"** – (*TreeBuilder*)
- **Related Coursework:** Operating Systems, Machine Learning Methods, High-Performance Computing Lab

### Technical University of Munich

Oct 2016 – Sept 2020

*B.Sc. Electrical and Computer Engineering*

*Munich, Germany*

## Publications

### Design and Implementation of A Scalable Financial Exchange in the Cloud | ([Link](#))

Jan 2024 – Present

- Novel Cloud financial exchange achieving low latency of  $\leq 250 \mu s$ , with a difference  $< 1 \mu s$  for 1K receivers.
- Achieves better scalability and around 50% lower latency than the multicast service provided by AWS.
- High-performance using kernel bypassing (DPDK) to reach a throughput of 35K multicast packet rate per second.

## Experience

### Research Assistant

Jul 2022 – Oct 2022

*TU Munich*

*Munich, Germany*

- Collaborated on **TensorDSE**, a Design-Space Exploration framework to accelerate machine learning model deployments.
- Assessed the performance metrics of multiple ML models across GPUs, CPUs and TPUs with TensorFlow Lite.
- Generated cost analysis reports on Google's Coral Edge TPU via USB traffic analysis (PyShark) during inference.
- Established that on average up to 57% of total inference time consists of data transmission with the external TPU.
- TensorDSE consumed reports to map a model's deployment optimally onto an available set of hardware devices.

### Embedded Software Engineer Intern

Aug 2021 – Jan 2022

*Molabo GmbH*

*Ottobrunn, Germany*

- Increased test coverage (GTest) up to 25% on safety-critical motor controller features, identifying and resolving bugs.
- Developed state simulation tooling with Linux's virtual CAN interface to validate motor functionality in real-time.
- Developed robust firmware update systems, ensuring reliability for partial updates (CAN bus) across 18+ clients.
- Actively developed the team's CI/CD build pipeline via Jenkins, Docker, and CMake, supporting over 10+ engineers.

### Tutor (Embedded Systems Programming Lab)

Apr 2021 – Aug 2021

*TU Munich*

*Munich, Germany*

- Mentored over 12 students on designing and developing low-level embedded FreeRTOS applications in C.
- Conducted 30+ sessions on best practices in data structures, concurrency, performance, and real-time scheduling.

## Technical Skills

**Languages:** C++, Rust, Python, Golang, Java, C, Bash, JavaScript, HTML, CSS, Lua, VHDL

**Cloud Services:** Google Cloud Platform (GCP), Amazon EC2 (AWS), Terraform, Docker, Packer, Vagrant

**Tools:** Linux, Unix Shell, Git, Github CI/CD, Jenkins, CMake, GNU Make, Bazel, Vim, VSCode, GDB, LLDB

**Technologies:** Docker, ZeroMQ, gRPC, Protobufs, MPI, TensorFlow, Scipy, NumPy, Pandas, DPDK

**Verbal/Written:** German – Fluent, Portuguese – Fluent

## Projects

### Cloud-TreeBuilder | *GCP, AWS, ZMQ, Terraform, Python, C++, Distributed Systems, Heuristic*

Mar 2024 – Present

- Optimally selects VMs in a cluster to form a multicast tree of depth D and fan-out F, to **minimize network latency**.
- Employs efficient **data pipelines** to gather VM network performance analytics (**UDP**) for informed heuristic selection.
- Improved **multicast latency by 24%** within the **microsecond** range of delivery for a cluster of 25 VMs.
- Integrated Terraform for cloud state management, ZMQ for node communication, and Scipy/Numpy for data analysis.

### Rust PKCS11 Client/Server | *Rust, Multithreading, HSM, Cryptography, Protobufs*

Dec 2024 – Present

- Implemented a PKCS11 **server** that accepts concurrent requests (**clients**) for typical cryptographic operations.
- Designed performant server backend with thread-safe access to SoftHSM for secure encryption, decryption and signing.

### Open-MPI Value Iteration | *C++, Parallel-Computing, MPI, HPC*

Mar 2022

- An HPC prototype that solves a stochastic navigation problem by distributing workload across an HPC cluster.
- Leveraged MPI to scale data distribution, achieving a 52% performance improvement over single-threaded execution.